
1. Problem

Autonomous AI systems fail in production when:

1. Identity drifts across LLM providers
2. Models invent policy instead of obeying governance
3. Inference cost and latency are uncontrolled
4. Memory is anecdotal rather than measured
5. Evidence is a screenshot, not a replayable artifact

The market does not need more answers. It needs governed work.

2. Architecture

Human / IDE

- > HIT (translation only)
- > HARLEMM (deterministic governance)
 - > Deterministic tasks
 - > InferenceRequired -> SULLY
- > ANNY (institutional identity)
- > Institutional Memory (measured vectors)
- > Evidence Factory (append-only, SHA-256)

Public product surfaces: Docs - SDK - Website - Whitepaper - Examples
Private institutional surfaces: Core - ANNY - HARLEMM - SULLY - Memory

3. Evidence First

Every institutional claim must be append-only, hash-addressed, replayable, DTO-linked, and time line-ordered.

CONRRAD does not invent green checks. Evidence precedes claim.

4. Institutional Memory

Canonical memory is measured in Qdrant collections with explicit writers, embedding contracts, latency, and recall.

Memory is infrastructure, not chat history folklore.

5. ANNY

ANNY is institutional identity. Across worker lanes:

GPT -> Cursor -> Antigravity -> Local -> GPT

structural identity remains continuous. The model may change; the institution must not.

6. HARLEMM

HARLEMM is a pure deterministic evaluator. It expands contracts, policies, and invariants and emits decisions.

It does not select LLMs and does not converse with humans.

7. SULLY

SULLY is the inference economy OS:

Difficulty -> Local RAM -> Local GPU -> Cursor -> Antigravity -> Response

GPU automation must not require human intervention when hardware exists.

8. Deterministic AI

CONRRAD separates:

- Deterministic work (HARLEMM / Core)
- InferenceRequired work (SULLY)

Determinism is the default. Inference is a controlled exception.

9. Economic Model

- Local SDK development is free
- agent.deploy() is the monetization boundary
- Infrastructure (Seal, Observatory, Economic Router) remains closed

10. YC Positioning

CONRRAD is infrastructure for governed autonomous AI ? the control plane for agents that must survive real work, audits, and cost pressure.

11. Competitive Advantage

| Others | CONRRAD |

|-----|-----|

| Chat wrappers | Institutional identity |

One-vendor lock-in	Model-agnostic routing
Unmeasured RAG	Audited institutional memory
Demo theater	Evidence-first product

Links

- Docs: <https://github.com/GRECOITALICO/CONRRAD-Docs>
- SDK: <https://github.com/GRECOITALICO/CONRRAD-SDK>
- Website: <https://github.com/GRECOITALICO/CONRRAD-Website>
- Examples: <https://github.com/GRECOITALICO/CONRRAD-Examples>